

Preemptive Rate-based Operator Scheduling in a Data Stream Management System *

Mohamed A. Sharaf, Panos K. Chrysanthis, Alexandros Labrinidis
Department of Computer Science
University of Pittsburgh
Pittsburgh, PA 15260, USA
{msharaf, panos, labrinid}@cs.pitt.edu

Abstract

Data Stream Management Systems are being developed to process continuous queries over multiple data streams. These continuous queries are typically used for monitoring purposes where the detection of an event might trigger a sequence of actions or the execution of a set of specified tasks. Such events are identified by tuples produced by a query and hence, it is important to produce the available portions of a query result as early as possible.

A core element for improving the interactive performance of a continuous query is the operator scheduler. An operator scheduler is particularly important when the processing requirements and the productivity of different streams are highly skewed. The need for an operator scheduler becomes even more crucial when tuples from different streams arrive asynchronously. To meet these needs, we are proposing a Preemptive Rate-based scheduling policy that handles the asynchronous nature of tuple arrival and the heterogeneity in the query plan. Experimental results show the significant improvements provided by our proposed policy.

1 Introduction

A *Data Stream Management System (DSMS)* hosts applications which rely on data that is continuously generated at remote sources. The DSMS is responsible for processing these data streams according to the applications' requirements and for streaming the relative results to each application. Such applications include monitoring a data network performance, fraud detection in telecommunication networks, monitoring the stock market, personalized Web

pages, and environmental monitoring via sensor networks.

Recently, a number of prototype system have been developed to accommodate the growing need for supporting the processing of data streams. Examples include: Aurora [3], STREAM [8], TelegraphCQ [5], and NiagaraCQ [6]. In these systems, continuous queries are registered at the DSMS and run continuously over the data streams. The arrival of new data triggers the execution of one or more continuous queries. The objective of these systems is to provide data processing techniques that is tailored for the continuous nature of data handled by DSMSs.

The requirements of a DSMS are more demanding than those of a traditional DBMS for several reasons. In a traditional DBMS, data is stored as relations and queries are posed on the stored relation. However, in a DSMS, data arrives in the form of continuous, unbounded, and possibly bursty data streams. Moreover, a DSMS typically supports multiple long-running *continuous queries* over the data streams [10].

Continuous queries are typically used for monitoring purposes where the detection of an event might require triggering a sequence of action. This creates an *interactive* environment where tuples produced by a query might activate the execution of specified tasks. In such an environment, it is important to optimize for producing the available portions of the result as early as possible rather than optimizing the computing of the entire result as in traditional database systems. The timely generation of the initial results increases the probability of detecting events earlier.

The two main elements for improving the interactive performance of a continuous query are: 1) *query optimizer*, and 2) *operator scheduler*. The query optimizer decides the dependencies between the query operators, whereas the scheduler decides the order of execution of the query operators. An operator scheduler is particularly important when a query is defined on multiple relations or streams.

In a naive approach, scheduling could be delegated to the

*This work is supported in part by NSF ITR Medium Award (ANI-0325353). The first author is supported in part by the Andrew Mellon Predoctoral Fellowship.

underlying operating system. However, this approach overlooks chances for improvement that are possible by exploiting the available information about the query operators. The Rate-based scheduling presented in [12] exploits such information for improving the interactive query performance in traditional databases and in Web databases. Basically, it employs a prioritizing scheme where the priority of an operator is defined in terms of its *cost* and *selectivity*.

Since the Rate-based policy was proposed for traditional database systems, it assumes priorities are static. However, in DSMSs, tuples continuously arrive at the system, some of which might correspond to a higher priority operator than the one currently executing. Thus, under Rate-based scheduling, high-priority operators might be blocked waiting for a low-priority one to complete its execution.

To alleviate this priority-inversion problem, in this paper, we propose a *Preemptive Rate-based* scheduling policy. The intuition is that preemption enables the adjustment of a previous scheduling decision based on the current state that is determined by the arrival of new data. The idea of using preemption in scheduling query operators has already showed success in reducing the memory requirements for processing bursty streams [1]. In this paper, we are also using the preemption mechanism but for improving the interactive performance of queries with stateless operators. Our experimental evaluation shows the significant improvements provided by our proposed scheduling policy compared to the existing ones.

The rest of this paper is organized as follows. In Section 2 provides an overview of query processing, optimization and operator scheduling. Our model and scheduling policies are presented in Section 3. Section 4 describes our simulation testbed, and then in Section 5 we discuss our experiments and results. We conclude in Section 6.

2 DSMS Background

2.1 Continuous Queries

A Data stream management system (DSMS), as shown in Figure 1, is designed to handle the processing of multiple queries over multiple data streams. In such a system, users register continuous queries that execute as new data arrives. Data arrives in the form of continuous streams. The arrival of new data is similar to an *insertion* operation in traditional database systems. A DSMS is typically connected to different data sources that generate streams at different rates. Moreover, a single query might receive its input from multiple streams (i.e., sources).

The query evaluation plan can be conceptualized as a data flow diagram [3, 1], which is a tree of nodes and edges, where the nodes are operators that process tuples and edges represent the flow of tuples from one operator to another

(Figure 2). An edge from operator O_1 to operator O_2 means that the output of operator O_1 is an input to operator O_2 . The overall query cost is determined by the cost of each operator C_i and the selectivity of each operator S_i . Recall that, an operator with selectivity S_i produces S_i tuples after processing one tuple for C_i time units. S_i is typically less than or equal to 1 for simple operators like filters, but, it might be more than 1 for joins.

A single continuous query in a DSMS could be quite complicated and expensive, especially, if the query is handling data from multiple streams where tuples from each stream pass through several expensive operators. Moreover, the distribution of the processing costs of operators is typically *heterogeneous*. That is, the costs for processing different tuples from different streams could be highly skewed. For instance, a query plan might contain operators that perform simple filtering on one data stream, whereas on another data stream complex operators might require join with stored relations or look-up indexes. Similarly, the selectivity of different operators exhibit a skewed distribution. This heterogeneity makes it essential for using query processing techniques that provide an efficient execution plan for evaluating the query.

2.2 Query Execution

In this work, we assume that each query is composed of only stateless operators. Queries with stateless operators is fairly common class of queries in data stream applications [1, 4]. Examples of stateless operators are *scan*, *selection*, *projection*, and *join* with stored relations. Moreover, a continuous query over multiple streams typically contains special operators for merging different streams. Such an operator is known as *union* [3], *merge* [7], or *mux* [9]. Studying queries with statefull operators (e.g., join over data streams) is part of our future work.

Every tuple that enters the system must pass through a unique sequence of operators, referred to as an *operator path* [1]. The operator path corresponding to stream Str might contain specialized operators that are only applicable on tuples from Str , in addition to other operators that are applicable on more than one stream. A specialized operator is basically used for transforming an input tuple to match a pre-defined scheme, or to apply a certain filtering or mapping for which the execution method depends on the corresponding data stream. For example, in Figure 2, operators forming the path O_1, O_2, O_3 process tuples from $Stream_1$, O_1 is applied only on tuples from $Stream_1$, while O_2 is used to merge the intermediate result tuples from $Stream_1$ and $Stream_2$. Finally, O_3 is used to project the final result tuples from both streams.

An operator receives its input tuples either directly from the input stream (e.g., operators O_1 and O_4) or from inter-

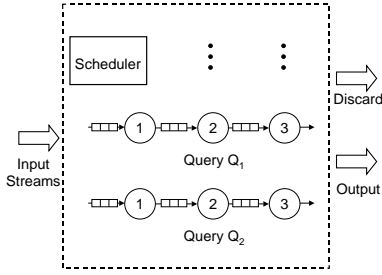


Figure 1. Data processing in DSMS

mediate queues used to buffer the output of a predecessor operator (e.g., operators O_2 and O_3). A tuple is processed until it is either produced at the output or discarded. A tuple is produced when it satisfies all the filters on the operator path. If a tuple does not satisfy a certain filter, it is immediately discarded since it is not used by any other query.

The mechanism of invoking an operator depends on the underlying query execution architecture. In the single-thread architecture, all operators are running in the same thread and invoking an operator is equivalent to a procedure call. In the multiple-thread architecture, each operator runs in a separate thread and invoking an operator requires context switching between threads.

The scheduler is the system component responsible for invoking operators according to a special order of execution. Specifically, it decides the following: 1) the order of query execution among the registered queries, and 2) the order of operator execution within a query. In this work, we are focusing on the latter problem of scheduling operators within a query. Specifically, we are interested into an operator scheduling policy that improves interactive performance of continuous queries.

2.3 Operator Scheduling

Traditionally, query optimizers are cost-based in that they decide among alternative execution plans by minimizing the query cost. That is, the estimated cost of evaluating the query until the last result tuple appears [13]. However, with the advent of the Internet, interactive query performance became an important criterion for evaluating the success of an online system. In such an interactive environment, it is important to optimize for producing the available portions of the result as early as possible rather than optimizing the computing of the entire result. The same interactive behavior is desirable in data stream systems for two main reasons. The first reason stems out of the nature of data stream processing; in DSMS data arrives continuously, hence, optimizing for producing the full result is obviously infeasible. The second reason depends on the timely requirements of the application supported by the DSMS. For




Figure 2. Query plan over streams

example, a user's continuous query might be used to process readings gathered by sensor networks that monitor environmental phenomena. In such a case, producing results as soon as they are available helps in accelerating the detection of abnormal behaviors that require immediate actions.

The two main elements for improving the interactive performance of a continuous query are: 1) *query optimizer*, and 2) *operator scheduler*. The query optimizer decides the dependencies between the query operators, whereas the scheduler decides the order of execution of the query operators. The work in [13] proposed a rate-based query optimization technique that maximizes the output rate of query evaluation plans. This technique provides optimization solutions for the case when a query contains operators that joins two streams which have different data arrival rates. For all other cases, the plans generated by the rate-based optimization are the same as those generated using a cost-based optimizer. Moreover, for both of the optimization strategies, the way operators are scheduled and how the flow of data is controlled can lead to significantly different kinds of output behavior for the same generated query plan [12].

For an optimized query plan, *Pipelined* query execution is one technique that can improve the interactive performance of a query. In a pipelined execution model, the unit of execution is the processing of a single tuple on a single stream. In a fully pipelined execution plan, result tuples are delivered as soon as they are computed from the input tuples. Pipelining works only as long as operators are non-blocking, i.e., they do not stage the data without producing results for a long time. Filters are non-blocking operators by nature, whereas for blocking operators (e.g., join), *windowing* is used to allow result streaming. However, for a multiple-relation query (or multiple streams), pipelined execution does not specify the order of execution of operators on different streams. A scheduling policy is needed to determine which stream to process when more than one has available tuples.

The problem of operator scheduling has been the focus of the work in [12]. It proposed a dynamic *rate-based pipeline scheduling* policy that produces more results during the early stages of query execution to improve a query

response time. Aurora [3, 4] also uses a technique similar to the rate-based pipeline scheduling to minimize the average tuple latency.

2.4 Performance Metric

The performance of scheduling policies is typically measured using the *average response time*. In traditional database systems, the response time of a query is defined as the amount of time from the time when the query is posed until the time when the last tuple in the result is produced.

However, in DSMSs, queries are continuous where a query is executed whenever new data arrives. That behavior is the opposite from traditional database systems; in traditional systems, the response is due to the arrival of a query, whereas in data streams system, the response is triggered by the arrival of a tuple.

Hence, in DSMS, it is more appropriate to define the response time from the data perspective rather than the query perspective. Therefore, we will define the *tuple response time* (or *tuple latency*) (R_i) for tuple i as follows:

Definition 1 $R_i = D_i - A_i$, where A_i is the tuple arrival time and D_i is the tuple departure time. Accordingly, the average response time for N tuples is: $\frac{1}{N} \sum_i^N R_i$.

Notice that tuples that are filtered out during query processing do not contribute to this metric [11].

Example Let us illustrate that measuring the tuple average response time better reflects the interactive performance achieved by a scheduling policy. In Figure 2, assume that the costs of operators O_1, O_2, O_3, O_4 are 10, 10, 10, 40 units respectively. Further, assume that there is one tuple available for processing at each input stream. First, consider a scheduling policy that will activate operators in the following order: O_1, O_2, O_3 to process the tuple from *Stream₁*, then O_4, O_2, O_3 to process the tuple from *Stream₂*. Such a policy will provide an average tuple response time of 60 units. An alternative scheduling might decide to schedule operator paths in the opposite order. That is, O_4, O_2, O_3 to process the tuple from *Stream₂*, then O_1, O_2, O_3 to process the tuple from *Stream₁*. This alternative policy will provide an average response time of 75 units. Though both policies take the same time to complete execution (i.e., 90 units), they provide different response times. The improvement exhibited by the first policy is due to producing result tuples as early as possible which is the behavior desired for interactive performance.

3 Operator Scheduling Policies

In this section, we first discuss two operator scheduling policies, namely, Round Robin and Rate-based, which are

currently used in the prototype DSMSs. Then we will describe our proposed preemptive version of the Rate-based scheduling policy.

3.1 Round Robin

Round Robin has been the policy used the most for processing sharing. It is simple and easy to implement. In Round Robin, each operator is assigned a time interval called quantum. At the end of the quantum, the processor is preempted and given to another operator.

By nature, Round Robin does not take advantage of the available parameters of operators (i.e., cost and selectivity). This is in contrast to using a priority-based scheduling policy which assigns each operator a priority based on its parameters. Ignoring this information makes Round Robin fall short in improving the interactive performance mentioned above. For example, it is known that a priority-based policy like *Shortest-Remaining-Processing-Time* significantly outperforms Round Robin in improving the response time [2]. However, Round Robin does not need to recover from a wrong scheduling decision because simply there is no decision taken and preemption allows all operators to execute sequentially.

3.2 Rate-based

The Rate-based scheduling policy was proposed for scheduling the execution of operators in a single query for the pipelined execution in traditional database systems [12]. Each operator O has a value called the *global output rate* which is defined in terms of its ancestor operators. The output rate of an operator O is basically the number of tuples produced per time unit by processing one tuple by the sequence of operators starting at O all the way up to the output. Formally, the global output rate for operator O_1 on the path O_1, O_2, \dots, O_n is computed as:

$$\frac{S_1 \times S_2 \times \dots \times S_n}{C_1 + C_2 \times S_1 + \dots + C_n \times S_{n-1} \times \dots \times S_1} \quad (1)$$

Where C_i and S_i are the cost and the selectivity of operator O_i , respectively.

The priority of each operator is set to its global output rate. At each scheduling point, the operator with the highest priority among all operators with non-empty input queues is the one scheduled for execution. In this policy, the scheduler is triggered (i.e., scheduling point is reached) when an operator finishes processing an input tuple. The new operator scheduled for execution can be the ancestor of the previously scheduled operator since it already has a higher priority. This higher priority is due to the increased probability of producing the tuple and/or the decrease in the remaining

execution time required to produce the tuple. Alternatively, the new scheduled operator might be the leaf node of a different operator path for which a corresponding tuple has arrived at the scheduling point or during the execution of the previous operator.

The intuition underlying this policy is to produce the initial output tuples as fast as possible. This is achieved by giving higher priority to operator paths that are productive and inexpensive or, in other words, select the operator path with the minimum latency for producing one tuple.

3.3 Preemptive Rate-Based

In data stream systems, tuples continuously arrive at the system; query execution is triggered by the arrival of new tuples. That is different from traditional database systems, for which the Rate-based policy was originally proposed. The problem with the continuous arrival of tuples is that a newly arriving tuple might correspond to a higher-priority path. Toward this, we are proposing a *preemptive* version of the rate-based policy so that a higher-priority path can be selected if such a path become available.

In this preemptive version, if operator O_i is already selected for execution then the next scheduling point could be triggered while O_i is still in progress. That is, scheduling occurs in the following cases:

1. **Operator O_i finishes processing:** This case is the same as in the non-preemptive version where a new scheduling decision is made every time an operator finishes executing.
2. **The arrival of a new tuple at a leaf operator O_j :** In this case, the priority of operator O_j is compared to the *current* priority of O_i where the latter is computed given the interval δ it already spent processing the current tuple. In other words, the global output rate of O_i is recomputed by replacing C_i with $C_i - \delta$ in Equation 1. Then the operator with the highest priority is scheduled for execution.

The intuition in using preemption is to be able to adjust the scheduling decision based on the current situation that is determined by the arrival of new data. In the absence of preemption, not only an inexpensive operator might have to wait until the currently relatively expensive operator finishes execution, but also, a productive operator might have to wait until a less productive operator finishes execution. Using preemption is particularly important when the distribution of operators costs and/or selectivities is highly skewed. In such a case, the priority assigned to operators span a long range of values. This is in contrast when the operators parameters are homogeneous, then the priority values assigned to operators are fairly close and blocking a

	Preemptive	Non-Preemptive
Pipelined	PRB	RB, FIFO
Non-Pipelined	RR	Train Processing

Table 1. Classification of Scheduling Policies

high priority operator until finishing the execution of a less priority one will have a little impact on the performance.

It is worth mentioning that if O_i is preempted due to the arrival of new a tuple at O_j , then O_i will not resume execution again until all the operators on the operator path starting at O_j finish execution. This is because of the following: 1) the priority of O_i does not increase while executing O_j , and 2) the priority of all the operators on the path from O_j to the output is higher than the current priority of O_i . The latter observation follows from the fact that O_j has a higher priority than O_i and that an ancestor of O_j will have a higher priority than O_j as explained above. Thus, a preempted operator is not considered again until higher priority operators finish execution. This shows the limited increase in scheduling overhead introduced by allowing preemption.

3.4 Discussion

Above, we described three policies for scheduling query operators, namely, Round Robin (RR), Rate-based (RB), and our proposed Preemptive Rate-based (PRB). In the next section, we provide a performance evaluation of these policies in addition to the *First-In-First-Out* (FIFO) scheduling policy. FIFO is implemented following the pipelined execution model where a tuple is completely processed along its operator path before the processing of a new tuple starts. However, in the presence of ready tuples from multiple streams, we use the tuple's *timestamp* as a tie-breaker to decide the order, where a tuple's timestamp is the time when the tuple entered the system.

Table 1 shows a classification of the four policies according to two features, namely, *preemption* and *pipelining*. The table also includes the *Train processing* mechanism used by Aurora for controlling the flow of tuples along a path [4]. A train is a sequence of tuples that are executed as a batch within an operator that is selected for execution using a rate-based policy. The goal of tuple train processing is to minimize the low-level overheads in a single-thread query execution architecture. In the case where the length of the train is one tuple, then train processing is fully pipelined and is similar to the Rate-based policy. Hence, we will only use the Rate-based policy as a representative of the two techniques.

It is worth mentioning that the proposed Preemptive Rate-based policy could be used for scheduling operators from multiple queries in order to improve the overall system response time. However, with multiple queries fairness

Parameter	Value
Query Depth	1 – 6
Query Fan-Out	2 – 50
Cost Skewness	0.0 – 0.9
Selectivity Skewness	0.0 – 0.9
System Utilization	0.9
Expensive Operator Depth	0 – 6
No. of Expensive Operators	0 – 32
Expensive Operator Cost	25 – 300 units
Overhead Cost	0 – 50 units

Table 2. Simulation Parameters

in scheduling is an issue to be considered. Such a problem is the focus of our future work.

4 Evaluation Testbed

In our experimental evaluation, we have used the average tuple response time metric from Section 2 to compare the query performance provided by the scheduling policies mentioned above. Accordingly, we simulated the execution of a continuous query as in [4] where the query is specified by two parameters: *depth* and *fan-out*. The depth of the query specifies the number of levels in the query tree and the fan-out specifies the number of children for each operator.

The costs and selectivities of the query operators are generated according to a Zipf distribution which is specified using the Zipf parameter. Setting the Zipf parameter to zero results in a uniform distribution, whereas increasing the value of the Zipf parameter increases the skewness of the generated values.

Each leaf operator is attached to a data stream where tuples arrive according to a Poisson distribution. The mean inter-arrival time for the Poisson distribution is set according to the simulated system utilization (or load). For a utilization of 1.0, the inter-arrival time is equal to the time required to complete the query execution, whereas for lower utilizations, the mean inter-arrival time is increased proportionally. All the results reported here are at a utilization of 0.9 and stream length of 10K tuples.

Finally, we use additional parameters to control the heterogeneity in the query plan, namely, *Expensive Operator Depth*, *No. of Expensive Operators*, and *Expensive Operator Cost*. Finally, we use the *Overhead Cost* parameter for calculating the overheads incurred by the different policies. The usage of the above parameters is explained along with the experiments in the next section. Table 2 summarizes all our simulation parameters.

5 Experiments

5.1 Skewness in Operators' Costs

In this experiment we set the query depth to 1 and the fan-out to 50. This is equivalent to a query that performs projection and/or filtering on 50 data streams then merge them together. The costs of the operators are generated in the range 1–100 with skewness 0.9 toward the inexpensive operators. The selectivity for all operator is set to 1.

Figure 3 shows the average tuple response time provided by the Round Robin (RR), First-In-First-Out (FIFO), Rate-based (RB) and the Preemptive Rate-based (PRB) policies. The figure shows that, in general, the average response time decreases by increasing the skewness of operators' costs. For instance, the minimum response time is achieved at a Zipf parameter of value 0.9 where most of the operators are inexpensive. In the case of Zipf parameter of value 0.0, the operators' costs are uniformly distributed which results in higher response time.

The figure also shows that as the degree of skewness increases, RR outperforms FIFO. Recall that Both RR and FIFO ignore the operators' characteristics (i.e., cost and selectivity), however, the figure shows that at high skewness, preemption can improve the performance. This observation is emphasized by comparing the response times provided by the Preemptive Rate-based (PRB) and the non-preemptive version (i.e., RB) where PRB always outperforms RB as well as RR and FIFO.

This improvement is better illustrated in Figure 4. Figure 4 shows the reduction in average response time provided by PRB compared to RB. The figure shows that the improvement increases by increasing the skewness. For example, the reduction is only 5% when the Zipf parameter is equal to 0.0 and it increases to 20% when the Zipf parameter is equal to 0.9. The increase in improvement is due to the increased heterogeneity in the query plan. In a plan with a highly skewed operators' costs, using RB might result in a case where a newly arriving tuple might have to wait for the currently executing tuple though the latter might have a relatively very high execution cost. PRB avoids this by allowing the preemption of the expensive operator.

5.2 Skewness in Operators' Selectivities

The setting for this experiment is similar to the previous one, however, we are changing the skewness of selectivity while setting the cost of all operators to 100 units. From now on, we will exclude RR and FIFO from the comparison due to their high response time as illustrated in Figure 3.

Operators take selectivities in the range 0.01 – 1.0 which are generated using Zipf. The skewness of the Zipf parameter is toward the low-selectivity operators. In this setting,




Figure 3. Cost Skewness

at a high value for the Zipf parameter, most of the streams will have low priority while very few will have a relatively high priority.

Figures 5 and 6 show the same behavior illustrated in the previous experiment. That is, PRB always outperforms RB and that the significance of reduction in response time increases by increasing the skewness. Hence, in order to avoid repetition, the remaining of this section will only report on experiments where we are varying the cost distribution.

5.3 Skewness' Position

In the previous experiments, we changed the skewness of the cost and selectivity distributions, yet, we had no control on the position of the skewness in the query plan. Figure 7 shows another experiment where we are controlling the location of skewness.

Specifically, we generated a query tree of depth 6 and fan-out 2 where we set the cost for all operators uniformly between 1 and 10 units. Then we introduced one *expensive operator* of cost 100. This expensive operator is located on the left-most operator path in the tree. However, the depth of the expensive operator is defined by a parameter *Expensive Operator Depth*. In the case where expensive operator depth is equal to 6, the operator at the bottom-left corner of the tree is the most expensive, accordingly, the left operator path of the tree is the only expensive path. Decreasing the depth of the expensive operator is equivalent to increasing the number of expensive path. In the case where expensive operator depth is equal to 0, the root of the tree is the most expensive operator, hence, all the query path are relatively expensive. Figure 7 shows that generally, the average response time decreases by increasing the expensive operator depth. This decrease is due to a decrease in the number of expensive paths as mentioned above.

In Figure 8 we are illustrating the reduction in response time provided by PRB compared to RB. Notice that the improvement increases by increasing the expensive operator depth (i.e., decreasing the number of expensive paths). This behavior is due to increasing the skewness in the query tree where preemption is needed. The improvement reaches a




Figure 4. Reductions by PRB vs. RB

maximum at expensive operator depth value of 4 (i.e., 4 expensive operators paths), then it starts decreasing for values of 5 and 6. At a value of 5, there are only 2 expensive paths in the tree and at a value of 6, there is only one expensive path. The low number of expensive paths decreases the chances of an expensive operator to be running at the arrival time of a new tuple that corresponds to an inexpensive operator.

5.4 Skewness' Amount

To further study the effect of skewness' location and amount, we conducted another experiment that is shown in Figure 9. The settings for this experiment is similar to the previous one, however, the location and the number of the expensive operators are set differently. In this experiment, we are varying the *number of expensive operators* between 1 and 32. Further, the expensive operators are all located at depth 6 (i.e., the leaf nodes). Hence, the number of expensive paths is always equal to the number of expensive operators.

Figure 9 shows the same behavior demonstrated in Figure 7. That is, the average response time increases by increasing the number of expensive operator path. Moreover, Figure 10 emphasizes the impact of the amount of skewness on the degree of improvement. For instance, at very low number of expensive operators/paths there are less chances of priority conflict between current and new tuples. As the number of expensive paths increases, the chances of a priority conflict increases. Further increase in the number of expensive paths brings the tree close to homogeneity, hence, less chances of priority conflict.

5.5 Skewness' Magnitude

Figure 11 shows the results of an experiment where we are changing the cost of the expensive operator. Specifically, there is one expensive operator in the query plan located at depth 4. The cost of this expensive operator takes the values between 25 and 300 time units. The figure shows that PRB outperforms RB. Moreover, it shows that RR can




Figure 5. Selectivity Skewness




Figure 7. Location of Skewness

also outperform RB when the cost of the expensive operator is set relatively high. This is because scheduling the expensive operator for execution might result in the arrival of many tuples that belong to inexpensive operator paths. In such a case, RR and PRB allow preempting the currently executing expensive operator and proceed with executing the less expensive one to produce more result tuples, whereas RB will wait for the expensive operator to finish then it will make new scheduling decisions. The higher the cost of the operator, the longer the RB will wait, and the more the increase in the average response time.

5.6 Scheduling Policy Overhead

Figure 12 shows the results of our last experiment where we are measuring the overhead of each scheduling policy. As mentioned earlier, the overheads depend on the underlying query execution architecture. However, instead of assuming a certain architecture, we counted the number of transitions between states during query execution. Transitions happen in the following cases: 1) invoking an operator, 2) preempting an operator, and 3) invoking the scheduler.

RR incurs the overheads of all three transitions: at the end of each quantum an operator is preempted and the scheduler is invoked to find the next operator in the cycle with available tuples which is then invoked for execution. In RB, the first and the third transitions take place, where when an operator finishes execution, the scheduler is invoked and the operator with the highest priority is executed. PRB incurs the same overheads as RB in addition to the overheads




Figure 6. Reductions by PRB vs. RB




Figure 8. Reductions by PRB vs. RB

due to enabling preemption. Specifically, the scheduler is invoked with the arrival of every new tuple which might result in preempting the currently executing operator.

In Figure 12 we assigned the same cost to each type of the transitions mentioned above. We call this parameter *overhead cost* and its takes the values between 0 and 60. The value of the overhead cost together with the number of transitions determine the total overhead incurred by a certain scheduling policy. In Figure 12, we measured the response time at the point where the expensive operator cost is 300 as in Figure 11. The figure shows the increase in response time with the increase in overhead cost. Moreover, it shows that RR outperforms RB only when the overhead cost is 0. As the overhead costs becomes 1, the performance of RR is highly unacceptable (note that in this figure, we use logarithmic scale for the Y-axis as opposed to the previous figures). This high overhead incurred by RR is due to the continuous preemption of operators. The rate of preemption is determined by the quantum length. In these experiments we assumed a quantum of length 1 time unit. Higher values for the quantum will show lower overhead, yet, significantly high compared to RB and PRB. The figure also shows that PRB scheduling outperforms the non-preemptive RB for values of overhead cost up to 50 time units. However, a value of 50 is an unrealistic overestimation of the overhead cost, since this represents an overhead equal to $\approx 16.6\%$ of the expensive operator cost. Recall that an expensive operator might require looking up indexes or performing a join with a stored relation, in either cases, the cost of that operator should be orders of magnitude the overhead cost.




Figure 9. Amount of Skewness




Figure 11. Magnitude of Skewness




Figure 10. Reductions by PRB vs. RB




Figure 12. Impact of Overheads

6 Conclusions

The rapid growth of DSMSs introduces more challenges to the database systems research. Challenges are mainly due to the continuous nature of data and the timely requirements of the continuous queries. This new environment requires rethinking the existing data processing techniques for efficient design of DSMSs. In this paper, we addressed one of those techniques, namely, query processing. Specifically, we focused on the operator scheduling component of query processing. The paper makes the following contributions:

1. It emphasizes the importance of the rate-based pipelined scheduling policies for scheduling query operators over multiple heterogeneous data streams.
2. It points out the importance of preemption in query operator scheduling which is particularly crucial to suit the asynchronous nature of tuple arrival.
3. It proposes the *Preemptive Rate-based* scheduling policy which combines the advantages of pipelined execution, rate-based scheduling, and preemption.

Our extensive experimental evaluation shows the significant improvements provided by our proposed Preemptive Rate-based policy. Currently, we are studying the problem of scheduling multiple queries over data streams.

References

- [1] B. Babcock, S. Babu, M. Datar, and R. Motwani. Chain: Operator scheduling for memory minimization in data stream systems. In *ACM SIGMOD Conf.*, 2003.
- [2] N. Bansal and M. Harchol-Balter. Analysis of SRPT scheduling: Investigating unfairness. In *ACM SIGMETRICS Conf.*, 2001.
- [3] D. Carney et al. Monitoring streams: A new class of data management applications. In *VLDB Conf.*, 2002.
- [4] D. Carney et al. Operator scheduling in a data stream manager. In *VLDB Conf.*, 2003.
- [5] S. Chandrasekaran et al. TelegraphCQ: Continuous dataflow processing for an uncertain world. In *CIDR Conf.*, 2003.
- [6] J. Chen, D. J. DeWitt, F. Tian, and Y. Wang. NiagaraCQ: A scalable continuous query system for internet databases. In *ACM SIGMOD Conf.*, 2000.
- [7] C. Cranor, T. Johnson, O. Spatschek, and V. Shkapenyuk. Gigascope: A stream database for network applications. In *ACM SIGMOD Conf.*, 2003.
- [8] R. Motwani et al. Query processing, resource management, and approximation in a data stream management system. In *CIDR Conf.*, 2003.
- [9] M. Sullivan. A stream database manager for network traffic analysis. In *VLDB Conf.*, 1996.
- [10] D. B. Terry, D. Goldberg, D. Nichols, and B. M. Oki. Continuous queries over append-only databases. In *ACM SIGMOD Conf.*, 1992.
- [11] F. Tian and D. J. DeWitt. Tuple routing strategies for distributed eddies. In *VLDB Conf.*, 2003.
- [12] T. Urhan and M. J. Franklin. Dynamic pipeline scheduling for improving interactive query performance. In *VLDB Conf.*, 2001.
- [13] S. D. Viglas and J. F. Naughton. Rate-based query optimization for streaming information sources. In *ACM SIGMOD Conf.*, 2002.